

Refinement of Reactive Systems

K. Morris¹, C. Snook², T.S. Hoang²,
G. Hulet¹, R. Armstrong¹, and M. Butler²

¹ Sandia National Laboratories, Livermore, California, U.S.A.

² ECS, University of Southampton, Southampton, United Kingdom

Abstract. Statechart notations with ‘run to completion’ semantics, are popular with engineers for designing controllers that react to environment events with a sequence of state transitions. However, they lack formal refinement and rigorous verification methods. Event-B, on the other hand, is based on refinement from an initial abstraction and is designed to make formal verification by automatic theorem provers feasible. We introduce a notion of refinement into a ‘run to completion’ statechart modelling notation, and leverage Event-B’s tool support for theorem proving. We describe the difficulties in translating ‘run to completion’ semantics into Event-B refinements and suggest a solution. We illustrate our approach and show how critical (e.g. safety) invariant properties can be verified by proof despite the reactive nature of the system. We also show how behavioural aspects of the system can be verified by testing the expected reactions using a temporal logic model checking approach.

Keywords: run-to-completion, statecharts, refinement

1 Introduction

Reactive Statecharts are open systems capable of receiving potentially non-deterministic input. This work exposes a shallow embedding of open Statecharts [11,12,13] semantics in Event-B. Statecharts provide a graphical language, generalized from state machines, that is popular with engineers, variants of which appear in Matlab Simulink/Stateflow [10] and the Ansys tools. Particularly attractive is providing accessibility to abstraction/refinement via Rodin/Event-B which has an intuitive metaphor in the Statechart semantics [11,12,13]. The commercial tools have similar ideas expressed as encapsulation and composition but not entailing any formal guarantees. The hope is that engineers can better understand the origin of proof obligations in refinements and achieve formal guarantees earlier in their designs where it is most tractable.

Previous work has developed a number of different semantics all with different purposes and outcomes [4,5,9]. Because this work is focused on a mapping to Event-B, safety property preserving refinement is key. Event-B provides not only a definition of refinement but a rubric for finding valid refinements and this is carried over into the Statecharts work presented here. In our version of Statechart semantics, refinement means a subsetting of traces from an abstraction. This has

the beneficial effect of preserving safety properties from abstraction to refinement and permits proofs to be discharged at the highest tractable level of abstraction. It is at the highest level of abstraction that proofs are presumably the easiest to discharge.

Many incompatible definitions of refinement have been posed by others [4,9] and that can lead to confusion. Though these separate refinements have different goals, all of which may be attractive to systems designers in different ways, preservation of safety proofs is the goal here. While these other forms of refinement cannot be said to conflict with the one presented here, they will not always preserve safety properties. From the Event-B vernacular it might be better to relabel these other approaches not as methods of model “refinement”, but rather methods of model “elaboration”.

Preservation of safety across refinement requires only a few restrictions to the original [5] Statecharts (e.g. transitions cannot cross containment boundaries arbitrarily), but still allows for both parallel and hierarchical composition. With these restrictions composition becomes a refinement, but not all refinements are compositions. Such a unification of composition and refinement can lead, not only to code reuse, but reuse of proofs.

Section 2 provides background material. Section 3 discusses the Statechart concept of ‘run to completion’ and how it can be specified in Event-B. Section 4 introduces our example case study; a drone. Section 5 gives an outline of our translation from State-Chart XML (SCXML) to Event-B. Section 6 discusses the types of refinement. Section 7 illustrates our approach to verifying safety invariant properties. Section 8 illustrates our approach to verifying control responses, and Section 9 concludes.

2 Background

2.1 SCXML

SCXML is a modelling language based on Harel statecharts with facilities for adding data elements that are modified by transition actions and used in conditions for their firing [17]. SCXML follows a ‘run to completion’ semantics, where trigger events³ may be needed to enable transitions. Trigger events are queued when they are raised, and then one is de-queued and consumed by firing all the transitions that it enables, followed by any (un-triggered) transitions that then become enabled due to the change of state caused by the initial transition firing. This is repeated until no transitions are enabled, and then the next trigger is de-queued and consumed. There are two kinds of triggers: internal triggers are raised by transitions and external triggers are raised by the environment (non-deterministically for the purpose of our analysis). An external trigger may only be consumed when the internal trigger queue has been emptied.

³ In SCXML the triggers are called ‘events’, however, we refer to them as ‘triggers’ to avoid confusion with Event-B

2.2 Event-B

Event-B [1,6] is a formal method for system design. It uses *refinement* to introduce system details gradually into the formal model. An Event-B model contains two parts: *contexts* and *machines*. Contexts contain *carrier sets*, *constants*, and *axioms* constraining the carrier sets and constants. Machines contain *variables* \mathbf{v} , *invariants* $I(\mathbf{v})$ constraining the variables, and *events*. An event consists of a guard denoting its enabled-condition and an action defining the value of variables after the event is executed. In general, an event \mathbf{e} has the form: **any \mathbf{t} where $G(\mathbf{t}, \mathbf{v})$ then $S(\mathbf{t}, \mathbf{v})$ end** where \mathbf{t} are the event parameters, $G(\mathbf{t}, \mathbf{v})$ is the guard of the event, and $S(\mathbf{t}, \mathbf{v})$ is the action of the event.

Machines can be refined by adding more details. Refinement can be done by extending the machine to include additional variables (*superposition refinement*) representing new features of the system, or by replacing some (abstract) variables by new (concrete) variables (*data refinement*). Event-B is supported by the Rodin Platform (Rodin⁴) [2].

2.3 UML-B State-machines

UML-B [14,15,16] provides a diagrammatic modelling notation for Event-B in the form of state-machines and class diagrams. The diagrammatic models relate to an Event-B machine and generate or contribute to parts of it. For example a state-machine will automatically generate the Event-B data elements (sets, constants, axioms, variables, and invariants) to implement the states. Transitions contribute further guards and actions representing their state change, to the events that they elaborate. State-machines are typically refined by adding nested state-machines to states.

Each state is encoded as a boolean variable and the current state is indicated by one of the boolean variables being set to **TRUE**. An invariant ensures that only one state is set to **TRUE** at a time. Events change the values of state variables to move the **TRUE** value according to the transitions in the state-machine.

While the UML-B translation deals with the basic data formalisation of state-machines it differs significantly from the semantics discussed in this manuscript. UML-B adopts Event-B's simple guarded action semantics and does not have a concept of triggers and run-to-completion. Here we make use of UML-B's state-machine translation but provide a completely different semantic by generating a behaviour into the underlying Event-B events that are linked to the generated UML-B transitions.

3 Run To Completion

The run to completion semantics is specified via an abstract basis that is extended by the model [12,13]. Figure 1 shows a statechart representation of how

⁴ An extensible toolkit which includes facilities for modelling, verifying the consistency of models using theorem proving and model checking techniques, and validating models with simulation-based approaches.

the basis enforces the run to completion semantics on the model transitions.

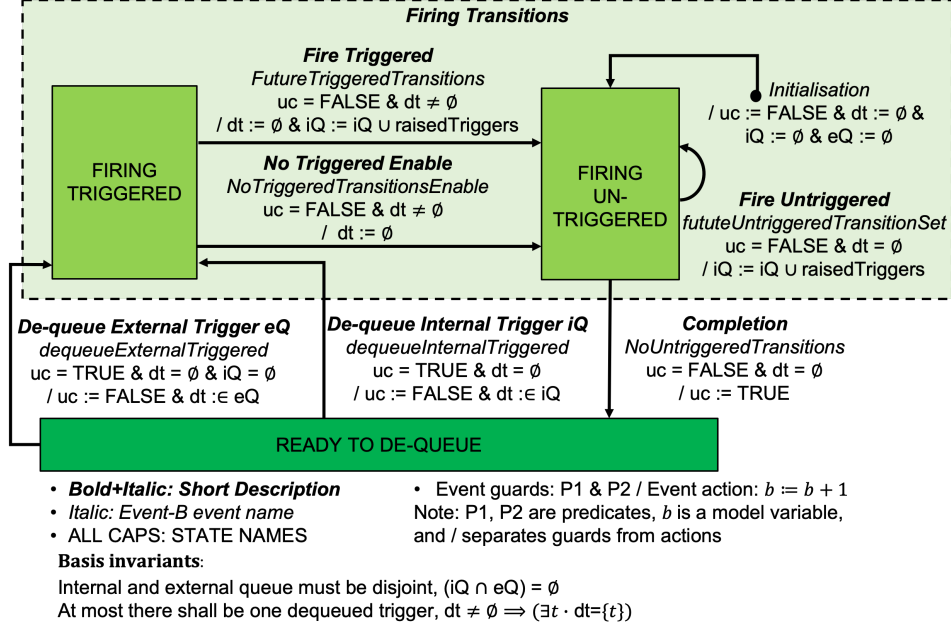


Fig. 1. Abstract representation of run to completion basis

The specification of this basis consists of an Event-B *context* and *machine* that are the same for all input models and are refined by the specific output of the translation. The basis context introduces a set of all possible triggers, **SCXML_TRIGGER**. This set is partitioned into internal and external triggers (e.g. **FutureInternalTrigger** and **FutureExternalTrigger** respectively), some of which will be introduced in future refinements. Each refinement partitions these trigger sets further to introduce concrete triggers, leaving a new abstract set to represent the remaining triggers yet to be introduced. For simplicity, this section describes the run to completion basis that is modelled using sets to abstractly represent the trigger queues. In Section 8 we will discuss the limitations of this abstraction and the ways in which they can be overcome.

Each of the transitions in the basis (see Figure 1) represents an abstract event of the basis machine that describes the generic behaviour of models under a run to completion semantics. These events provide an abstraction that defines the altering of trigger queues and completion flag. In order to preserve refinement rules, which prohibit the modification of existing variables, all future events generated by translation of the specific SCXML model, must refine the abstract events of the basis. The basis machine also declares variables that correspond to the triggers ready to dequeue at any given time, dt , the queue with internal triggers raised by actions within the model, iQ , the external triggers queue with triggers raised by the environment, eQ , and a flag, uc , that signals when a run

to completion macro-step has been completed (no un-triggered transitions are enabled).

In the **Initialisation** event, all trigger queues are emptied and **uc** is set to **FALSE** so that un-triggered transitions are dealt with via the **futureUntriggeredTransitionSet** event, which may also raise internal triggers. This will subsequently enable completion and reset the **uc** flag to **TRUE**. The abstract event **futureRaiseExternalTrigger** represents the raising of an external trigger (this transition is not shown in the diagram). After completion, raised triggers are dealt with by moving them to the dequeuing queue, **dt**. Internal triggers have a higher priority, since the external trigger queue is only dequeued if the **iQ** is empty (see **dequeueExternalTriggered** and **dequeueInternalTriggered** in Figure 1). The abstract event **futureTriggeredTransitions** represents a combination of transitions that are triggered by the trigger presently ready to dequeue, **dt**. The actions of these transitions may also raise triggers of their own.

In the process of refining a model, a designer takes advantage of the non-determinism in the abstraction. When a refinement level for which the designer wants to enforce a requirement is reached, the model needs to be *finalized* (see Section 5 for more on finalisation). The SCXML translation tool will then automatically strengthened the guards of events **NoTriggerTransitionEnable** and **futureUntriggeredTransitionSet**, to ensure that the run to completion sequence is not interrupted by non-deterministic behaviour. The guards of event **NoTriggerTransitionEnable** are strengthened by adding the conjunction of the negated guards of all transitions triggered by **dt**. In the case of the **futureUntriggeredTransitionSet** event, guard strengthening is achieved by adding the conjunction of the negated guards of all untriggered transitions.

4 Description of the Sample Application

To illustrate the development and analysis process of a design using the previously described statechart semantics, we will discuss a quadrotor helicopter or quadrotor application similar to the one presented by Syriani et al. [4]. The application will focus on the incremental design of some of the drone's required functionality. The constructed model must obey statechart refinement rules that are proven within the Rodin tool. The structure of the statechart for this model at each subsequent abstraction level restricts further the development of the model to refinements that obey the rules. This will allow us to prove properties of the model in a very strategic fashion, as properties proven of early abstraction levels are preserved in later refinements.

The first abstraction of the model shown in Figure 2 captures the basic functionality of the drone. The model's initial state is **OFF** and as a result of the **on** and **toTakeoff** external triggers it transitions to the **START** and **OPERATIONAL** states respectively⁵. The drone reacts to the **off** external trigger by shutting down and subsequently transitioning to the **OFF** state. Within the **OPERATIONAL**

⁵ Transitions in Figures 2–3 are labeled with trigger names (e.g. toTakeoff, toFly) not with event names as it is in UML-B.

state the drone will transition to **FLY**, **DESCEND** or **LANDED** state after the internal trigger **toFly**, **toLand** or **landed** is raised, respectively. In this abstraction, these internal triggers are raised non-deterministically in the system by functionality not currently defined. As additional details are incorporated into the model in later refinements some of that non-determinism is removed and replaced by transitions with actions that raised the previously defined internal triggers. It should be noted that this abstraction of the drone model includes a transition from **TAKEOFF** to **DESCEND** (dashed transition in Figure 2). This allows for the drone to respond to a **toLand** trigger if it encounters some problems while in the **TAKEOFF** state. Syriani et al. [4] introduces this transition in later refinements under Rule 8 *path refinement rule*. This rule is inconsistent with our rules of refinement as it results in a concrete event with no corresponding behavior in the abstraction.

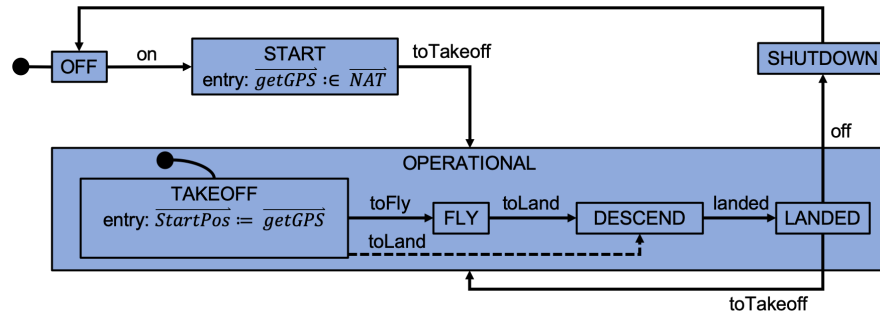


Fig. 2. Statechart of drone application. Abstract level including only generic behavior.

Removing one of the figures to reduce page count. Updates to next paragraph to agree with figure.

Figure 3 shows three subsequent refinements to the drone model. The first refinement of the model is shown in beige, as we refine the parent state **TAKEOFF** by introducing child states and new model variables, similar to Rule 2 *basic-to-or state rule* defined by Syriani et al. [4]. As part of this refinement we introduced an untriggered transition responsible for raising the **toFly** internal trigger, and therefore removed some of the non-determinisms in the abstraction.

The second refinement, the details of which are shown in green in Figure 3, extends the capabilities within **OPERATIONAL** by making it a parallel state that controls flying and battery related functionality. This is the same as Rule 4 *and-state rule* defined by Syriani et al. [4]. The charge within the drone battery is control by the parallel **BATTERYOP** state. The functionality is modeled by introducing a new model variable, **charge**, which is decreased as a response to the internal trigger **decreaseCharge**. The aforementioned trigger, is raised non-deterministically by some unspecified internal functionality. Our statechart semantics supports transition refinement, as such we are able to modified previously defined transitions. In particular, this type of refinement allow us to add guards and/or actions to previously defined transitions. The strengthening of guards, or additional actions are expressed in term of new model variables that contribute implementation details to the model. To ensure the drone operates

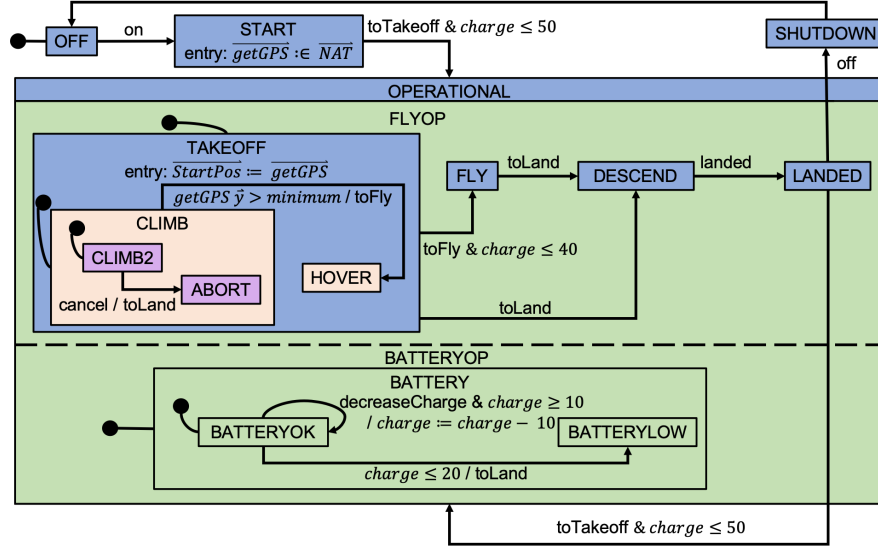


Fig. 3. Statechart of drone application. Refinement level for battery consumption functionality (shown in green). Refinement level for descending capabilities, in case of emergency (shown in lilac).

with enough battery power we strengthen the guards of transitions to the **FLY** and **TAKEOFF** states. As part of this design stage we introduce a requirement to constrain drone operation to a battery charge of at least 20% capacity. This can be expressed as

$$(\text{BATTERYOK} = \text{TRUE}) \Rightarrow \text{charge} > 20\% .$$

Figure 3 shows the third refinement of the drone model, with features added in lilac. At this stage we introduce additional implementation details to ensure that under special circumstances (e.g. sensing of adverse environment or unexpected battery dropped) the drone is able to circumvent flying and proceed to an emergency landing. The previously described requirement can be expressed as

$$(\text{TAKEOFF} = \text{TRUE}) \Rightarrow (\text{BATTERYOK} = \text{TRUE} \vee \text{toLand}) .$$

To implement this new capability in the design the internal trigger **cancel** is introduced. The internal trigger **cancel** can be raised non-deterministically by some sensing capability, the details of which are not currently implemented. If the trigger is raised, the climbing process must be aborted and the drone descending sequence shall start. This refinement level is done differently from Syriani et al. [4], which follows Rule 7 *state extension rule*. The aforementioned rule requires a data remapping of the abstract states **TAKEOFF**, **CLIMB** and **HOVER**, which should be distinct from the states in this refinement, as the state **ABORT** is introduced. In contrast, we implement this refinement using a rule

similar to Syriani et al.'s Rule 2 *basic-to-or state rule*, which introduces the concrete states **CLIMB2** and **ABORT** to the abstract state **CLIMB**.

5 SCXML Translation

The translation of a specific SCXML model comprises the following stages:

- Firstly, a basis machine and context are created to embody the semantics of the SCXML language (Section 3). The basis provides variables and events to model the queue of triggers as well as abstract versions of events to model transitions firing. The basis is independent of the particular SCXML model which is added in subsequent refinements.
- Secondly, all possible combinations of transitions that can fire together are calculated and corresponding events are generated, at appropriate refinement levels, that refine the abstract basis events. If these transitions raise internal triggers, a guard, (e.g. $\{i1, i2, \dots\} \subseteq \text{raisedTrigger}$, where $i1, i2, \dots$ have been added to the internal triggers set), is introduced to define the raised triggers parameter. The subset allows more triggers to be raised in later refinements. For triggered transitions, the trigger is specified by a guard that defines the value of the trigger parameter.
- Thirdly, the SCXML state-chart is translated into a corresponding UML-B state-machine whose transitions elaborate (i.e. add state change details to) the transition combination events that the transition may be involved in. A transition may fire in parallel with transitions of parallel nested state-machines that have the same (possibly null) trigger.
- Finally the UML-B state-machine is translated into Event-B by programmatically invoking the UML-B translator.

Deleted section listing all new feature in the translator since this is the ABZ paper.

Further details of the translation are given in [11,12,13].

A tool to automatically translate SCXML models into UML-B has been produced. The tool is based on the Eclipse Modelling Framework (EMF) and uses an SCXML meta-model provided by Sirius [3] which has good support for extensibility. The UML-B state-machine is subsequently translated into Event-B using the standard UML-B translation [15] which provides variables to model the current state and guards and actions to model the state changes that transitions perform.

6 Statechart Refinement

Our system includes three refinement rules.

1. Guard conditions on a transition can be strengthened; this can be done by adding textual guards to the transition, or changing the source of the transition to a nested state.

2. Transitions can have additional actions, provided they do not modify variables appearing in the abstraction; this can be accomplished by adding textual action to the transition or by changing the target to nested state.
3. A statechart can be embedded within a state of another statechart – sometimes called hierarchical composition or hierarchical refinement.

Via the translation explained in Section 5, these rules rely on the usual Event-B proof obligations to ensure that they do indeed yield refinements in the Event-B semantics.

If an Event-B model **B** can be shown (via the construction rules of the Event-B language as well as the proof obligations) to refine another Event-B model **A**, then we know that every behavior of **B** is also a behavior of **A**. This definition yields a useful principle of preservation of safety – if we can show that a bad thing never happens in **A**, then we can add detail via refinements in **B**, knowing that the bad thing will continue to never happen in **B**. That is, Event-B refinements preserve safety properties in the sense of [8]. This makes refinement a useful technique in developing safety-critical systems: one can analyze a simpler abstract model for critical safety properties and then add detail to the model via refinements, secure in the knowledge that the safety properties will be preserved.

Event-B refinements have also been shown to preserve some liveness properties, under certain conditions [7].

Although the autonomous drone example in this paper is based on the example described in [4], the definition of refinement used in that work is quite different from our own. This forces some differences in our refinement rules and consequently the way the example is developed. For example, the transition triggered by **toLand** (see 3 transition from **TAKEOFF** to **DESCEND**) adds new behaviour. In [4] “refinement” is a transformation of the model which preserves reachability of a state with respect to sequences of inputs. This guarantees that a refinement in their system will include all the behaviors of the abstraction, while possibly adding others. While this notion of refinement seems useful in certain contexts, unlike refinement in Event-B it does not guarantee preservation of safety properties. Therefore it should be considered less suited to development of safety-critical systems.

7 Verification of Safety Properties

In a state-chart model we naturally wish to verify properties **P** that are expected to hold true in a particular state **S**. Hence, all of the safety properties that we consider are of the form: $S = \text{TRUE} \Rightarrow P$, where the antecedent is implicit from the containment of **P** within **S**.

There are two kinds of properties that we might want to verify in an SCXML state-chart; 1) properties concerning the values of auxiliary data maintained by the system and 2) constraints about the state of another parallel state-chart region.

SCXML models represent components that react to received triggers and cannot be perfectly synchronised with changes to the monitored properties. Hence,

P may be temporarily violated until the system reacts by leaving the state S in which the property is expected to hold. To cater for this we express P in a modified form P' that allows time for the reaction to take place.

There are two forms of reaction that can be used to exit S ; a) an untriggered transition, or b) a transition that is triggered by an internally raised trigger. For a), the modified property P' becomes $P \vee \text{untriggered transitions are not complete}$, and for b) P' becomes $P \vee \text{trigger } t \text{ is in the internal queue or dequeued}$ (where t is the internal trigger raised when the violation of P is detected).

In this section we illustrate a typical example of the type of properties that we imagine could be verified in a reactive SCXML system. All of the proof obligations are automatically discharged for our example. Since our models are strictly structured and proof obligations will always have this common form, we are optimistic that proofs will always discharge automatically. We model the safety property features at an early level of refinement where the models are relatively simple, so that the validity of verification conditions is clear. Detail is then added in later refinements which are proven (automatically) to preserve the previously verified safety properties. In our example, some auxiliary data is monitored by one state-chart region and while a parallel region refers to the state of the monitoring region. Hence the reaction consists of an un-triggered transition in the monitoring region which sends an internal trigger to the other region when it leaves the desired monitor state.

For our drone model, the safety property that we wish to verify is that the control system does not continue to take off or fly if the battery charge drops below a certain threshold (say 21%). By refinement level 1 we have developed the drone's state to the point where we distinguish the **TAKEOFF** and **FLY** states (Figure 2). In refinement level 2 we therefore introduce the battery charge monitoring function along with the associated safety properties. A parallel state-chart region, with sub-states **BATTERYOK** and **BATTERYLOW**, is added to the state **OPERATIONAL** (Figure 3). The **BATTERYOK** sub-state is used in the safety invariant of the **TAKEOFF** and **FLY** states. Thus we split the verification into two parts: a *type b* proof to show that the system reacts to the battery charge decreasing below 21% (an external event) by leaving the **BATTERYOK** sub-state, and a *type a* proof to show that when the system leaves the **BATTERYOK** state it subsequently (within the run to completion) leaves the **FLY** or **TAKEOFF** states. Both parts are described in more detail as follows.

System Reacts to the Low Battery Charge An external trigger indicates that the battery charge has dropped by 10% and this is used by a self transition to decrement the controllers data value for charge. The **BATTERYOK** state is supposed to indicate that the battery charge is ok ($>20\%$) and to ensure that it does, we add a state invariant to this effect ($\text{charge} > 20$). When charge decreases to 20 (or less), an untriggered transition immediately reacts by switching to the **BATTERYLOW** state. To ensure that this reaction is not bypassed by the non-determinism that we incorporated to allow for future refinement, we flag it as finalised at refinement level 2. Finalisation means that we cannot strengthen its guards in future refinements as is normally permitted, since its reaction is

needed to ensure the invariant is preserved. After translation to Event-B via UML-B the invariant in state **BATTERYOK** is

$$(\text{BATTERYOK} = \text{TRUE}) \Rightarrow (\text{uc} = \text{FALSE} \vee \text{charge} > 20) .$$

The only events that can break this invariant are ones that make the antecedent become true or the consequent become false and we deal with these as follows: The transitions that enter state **OPERATIONAL** and initialise the **BATTERY** region by entering **BATTERYOK** (hence making the antecedent become true) contain the guard that $\text{charge} > 50$ (since we do not allow the drone to take off unless the battery is well charged) and hence the invariant is satisfied. The self transition that decreases charge (and hence could potentially falsify the consequent) is guarded by $\text{uc} = \text{FALSE}$ since it is a triggered transition, and hence the disjunction in the consequent ensures it remains true. The completion event **NoUntriggeredTransitions** of the basis machine resets $\text{uc} = \text{TRUE}$ to indicate completion of the cycle and hence could potentially break the invariant. However, finalising the transition **BATTERYOK.BATTERYLOW** (that leaves **BATTERYOK** when $\text{charge} > 20$ becomes false) means that the negation of its guard is added to the completion event by the translation. Since this transition fires when $\text{BATTERYOK} = \text{TRUE}$ (i.e. its source state) and $\text{charge} \leq 20$ the completion event is guarded by $\neg(\text{BATTERYOK} = \text{TRUE} \wedge \text{charge} \leq 20)$ which means that it does not fire when it could break the invariant (i.e. forcing the untriggered reaction to fire first).

System Subsequently Leaves the FLY or TAKEOFF States The safety property of the **TAKEOFF** and **FLY** states can now be simply stated as $\text{BATTERYOK} = \text{TRUE}$. However, since this relies on a particular internal trigger (**toLand**) to make the appropriate reaction, we also need to specify that trigger as an attribute of the invariant in the SCXML model. After translation to Event-B via UML-B the invariant in state **TAKEOFF** becomes

$$(\text{TAKEOFF} = \text{TRUE}) \Rightarrow (\text{toLand} \in \text{iQ} \vee \text{toLand} \in \text{dt} \vee \text{BATTERYOK} = \text{TRUE}).$$

The invariant for the **FLY** state is similar with a corresponding antecedent. The transitions that enter **TAKEOFF** (which make the antecedent true) simultaneously enter **BATTERYOK** ensure the consequent is true. The only transition that enters **FLY** (which makes the antecedent of the **FLY** invariant true) comes from the **TAKEOFF** state and hence the consequent is already true. The transition that leaves **BATTERYOK** (making the last disjunct of the consequent false) raises the **toLand** trigger making the first disjunct true. Some transitions leave the superstates of **BATTERYOK** but these either simultaneously leave **OPERATIONAL** (the superstate of **TAKEOFF** and **FLY**), or re-enter **BATTERYOK**. The basis contains an event to dequeue the internal triggers which preserves the overall consequent because establishes the second conjunct as it falsifies the first (i.e. it removes **toLand** from the **iQ** but simultaneously adds it to **dt**). The only events that falsify the second conjunct are the transitions triggered by **toLand** which leave the **TAKEOFF** or **FLY** states making the antecedent false.

Hence, invariant properties that follow these suggested patterns are always automatically proven due to simple logic about the changes in state.

8 Verification of Control Responses

It is sometimes possible to construct a model that satisfies some invariant (e.g. safety) properties, but does not behave in a useful way. Therefore, as well as verifying invariant properties, we would like to verify the system's responsiveness. More specifically in this case, we want to ensure that the controller responds to external triggers to make appropriate modifications to the system variables. These kind of live responses can not be verified by proof of invariants since they are temporal properties. Instead, we can express the property in Linear Temporal Logic (LTL) and use the ProB⁶ model checker to verify it.

In general, our liveness properties will have the following form:

$$G([\text{external_trigger_event}] \Rightarrow F\{\text{predicate}\}) ,$$

where the predicate concerns variables v that the system maintains, and may refer to old values $\text{old}(v)$ that existed when the external trigger occurred. To specify a liveness property to be verified, a special LTL element is added to the SCXML model with attributes, property (a string of the above form) and refinement (an integer indicating the refinement level at which the property should be verified). The translator generates a separate 'branch' refinement for each LTL property to be verified. In this special refinement, history variables are added to record the value at the state when the external trigger occurs, of any variables that are referenced as 'old' values. A text file is automatically generated containing the LTL property to be checked. In this generated version, an assumption of strong fairness is added for all other events in the model. (This assumption is stronger than necessary since some events will not affect the outcome, but is easier to generate and is sufficient for our verification aim). For simplicity we omit this assumption from the remaining examples.

$$SF[e1] \wedge SF[e2] \dots \Rightarrow G([\text{external_trigger_event}] \Rightarrow F[\text{predicate}])$$

~~This property can be copied and pasted into the ProB model checker LTL checking text field.~~ This property can be added into the ProB model checker LTL formula text field.

We illustrate the method with an example of a temporal property that we expect to hold in the drone SCXML system. The liveness property that we wish to verify is that, after an external trigger event `decreaseCharge`, the battery charge value should decrease in value.

$$G([\text{ExternalTriggerEvent.decreaseCharge}] \Rightarrow F\{\text{charge} < \text{old}(\text{charge})\}) .$$

At first, we could not verify this property. The counter example traces that ProB provided gave us a better understanding of the reasons why:

⁶ ProB is an animator, constraint solver and model checker for the B-Method.
<https://www3.hhu.de/stups/prob>

- Our model represented the trigger queues abstractly as sets which meant that the **decreaseCharge** trigger may never be dequeued. The standalone version of ProB allows strong fairness to be specified for particular parameter values but this does not work in the Rodin plug-in for ProB. In any case, a more accurate (concrete) representation of the queue fixes the problem and improves our model.
- The charge is not always decreased in response to the **decreaseCharge** trigger. The controller only monitors battery charge while in the **BATTERYOK** state and discards the trigger in other states. Also, the controller stops decreasing charge when it approaches 0. To cater for this we added a pre-condition $\text{BATTERYOK} = \text{TRUE} \wedge \text{charge} \geq 10$ to the LTL property.
- Even if this pre-condition is true when the trigger is raised, another trigger (e.g. **off**) may already be in the queue and take the controller out of **BATTERYOK** before the **decreaseCharge** trigger is dequeued. Again we strengthen the pre-condition $\text{off} \notin \text{dt} \cup \text{eQ}$ of the LTL expression to avoid this situation.

After making these changes the final form of the LTL property, which ProB was able to exhaustively check and confirm was as follows:

$$G([\text{ExternalTriggerEvent.decreaseCharge}] \wedge \{\text{BATTERYOK} = \text{TRUE} \wedge \text{charge} \geq 10 \wedge \text{off} \notin \text{SCXML_dt} \cup \text{SCXML_eq}\} \Rightarrow F \{\text{charge} < \text{old}(\text{charge})\}) .$$

9 Conclusion

Reactive Statecharts are useful and widely used by engineers for modelling the design of control systems. Event-B provides an effective language for formally verifying properties via incremental refinements. However, it is not straightforward to apply the latter to the former. We have demonstrated a technique for introducing refinement of reactive Statecharts that can be translated to Event-B for verification. Invariant properties about the expected coordination of states can be added and are interpreted with additional allowance for the reactions to take place. That is, they hold only after the reaction has taken place. Such invariants prove automatically with the existing Rodin theorem provers. We also demonstrate a complementary process for verifying expected reactions to environmental triggers that uses the LTL model checker. In future work we intend to formalise the semantics of our extended SCXML notation in order to define its notion of refinement and correspondence to Event-B.

All data supporting this study are openly available from the University of Southampton repository at <http://doi.org/10.5258/SOTON/D0??> (Temporary link: <https://tinyurl.com/droneModel>)

References

1. Abrial, J.R.: Modeling in Event-B: System and Software Engineering. Cambridge University Press (2010)

2. Abrial, J.R., Butler, M., Hallerstede, S., Hoang, T., Mehta, F., Voisin, L.: Rodin: An open toolset for modelling and reasoning in Event-B. *Software Tools for Technology Transfer* **12**(6), 447–466 (Nov 2010)
3. Eclipse Foundation: Sirius project website.
<https://eclipse.org/sirius/overview.html> (Mar 2016)
4. Eugene Syriani, Vasco Sousa, L.L.: Structure and behavior preserving statecharts refinements. *Science of Computer Programming* **170**(15), 45–79 (Jan 2019), <https://doi.org/10.1016/j.scico.2018.10.005>
5. Harel, D.: Statecharts: A visual formalism for complex systems. *Sci. Comput. Program.* **8**(3), 231–274 (Jun 1987). [https://doi.org/10.1016/0167-6423\(87\)90035-9](https://doi.org/10.1016/0167-6423(87)90035-9), [http://dx.doi.org/10.1016/0167-6423\(87\)90035-9](http://dx.doi.org/10.1016/0167-6423(87)90035-9)
6. Hoang, T.S.: An introduction to the Event-B modelling method. In: *Industrial Deployment of System Engineering Methods*, pp. 211–236. Springer-Verlag (2013)
7. Hoang, T.S., Schneider, S., Treharne, H., Williams, D.: Foundations for using linear temporal logic in event-b refinement. *Formal Aspects of Computing* **28** (04 2016). <https://doi.org/10.1007/s00165-016-0376-0>
8. Lamport, L.: Proving the correctness of multiprocess programs. *IEEE Transactions on Software Engineering SE-3* **2**, 125–143 (March 1977), <https://www.microsoft.com/en-us/research/publication/proving-correctness-multiprocess-programs/>
9. Maraninchi, F.: The Argos language: Graphical representation of automata and description of reactive systems. In: *IEEE Workshop on Visual Languages* (1991)
10. MATLAB: 9.7.0.1190202 (R2019b). The MathWorks Inc., Natick, Massachusetts (2019)
11. Morris, K., Snook, C.: Reconciling SCXML statechart representations and Event-B lower level semantics. In: *HCCV - Workshop on High-Consequence Control Verification* (2016), http://www.sandia.gov/hccv/_assets/documents/HCCV_2016_Morris.pdf
12. Morris, K., Snook, C., Hoang, T.S., Armstrong, R., Butler, M.: Refinement of statecharts with run-to-completion semantics. In: Artho, C., Ölveczky, P.C. (eds.) *Formal Techniques for Safety-Critical Systems*. pp. 121–138. Springer International Publishing, Cham (2019)
13. Morris, K., Snook, C., Hoang, T.S., Hulette, G., Armstrong, R., Butler, M.: Refinement and verification of responsive control systems. In: Raschke, A., Méry, D., Houdek, F. (eds.) *Rigorous State-Based Methods*. pp. 272–277. Springer International Publishing, Cham (2020)
14. Said, M.Y., Butler, M., Snook, C.: A method of refinement in UML-B. *Softw. Syst. Model.* **14**(4), 1557–1580 (Oct 2015). <https://doi.org/10.1007/s10270-013-0391-z>, <http://dx.doi.org/10.1007/s10270-013-0391-z>
15. Snook, C.: iUML-B statemachines. In: *Proceedings of the Rodin Workshop 2014*. pp. 29–30. Toulouse, France (2014), <http://eprints.soton.ac.uk/365301/>
16. Snook, C., Butler, M.: UML-B: Formal modeling and design aided by UML. *ACM Trans. Softw. Eng. Methodol.* **15**(1), 92–122 (Jan 2006). <https://doi.org/10.1145/1125808.1125811>, <http://doi.acm.org/10.1145/1125808.1125811>
17. W3C: State chart XML SCXML: State machine notation for control abstraction. <http://www.w3.org/TR/scxml/> (Sep 2015)